

Data Cleaning with R



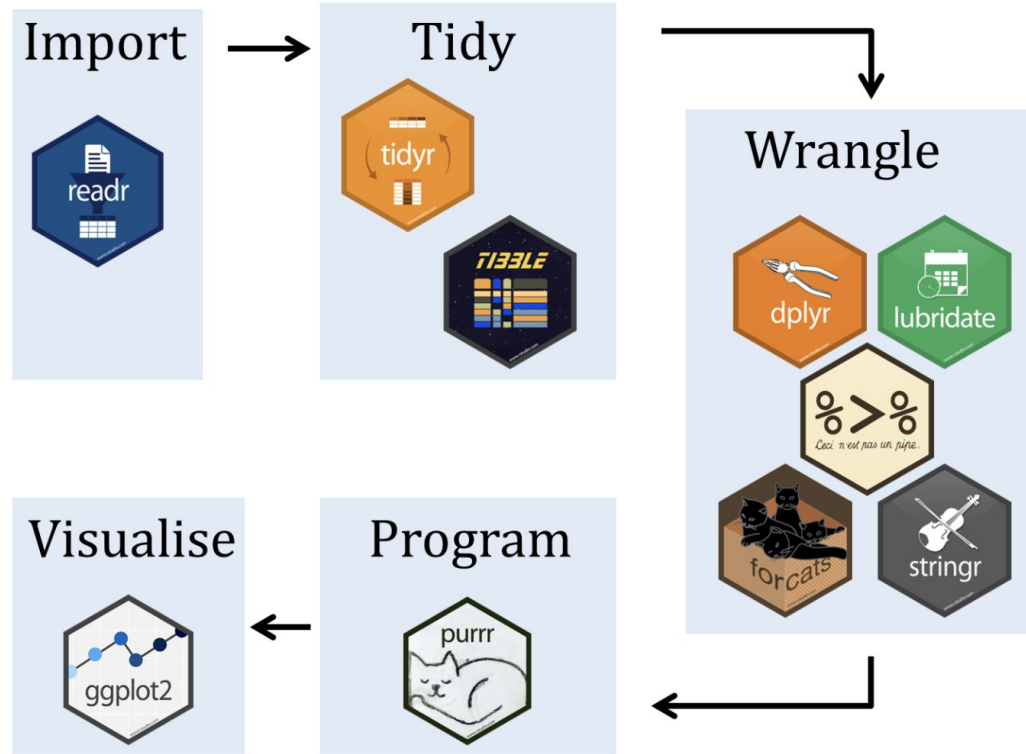
Messy data can be caused by

- Error in data entry
- Incomplete data entry
- Bringing data in from multiple sources or streams
- Changes in how the data is organized over time
- Problems introduced when data is exported/shared/opened

What data cleaning tasks are common?

- Fix errors in the data
- Removing duplicate records
- Removing outliers
- Stripping whitespace, removing unwanted characters, changing case
- Restructuring the data
- Splitting data into multiple columns
- Merging data

tidyverse: R packages for data cleaning and visualization



tidyverse packages expect *tidy* data

(packages tidyr and tibble are designed to help you get the data in tidy format)

country	year	cases	population
Afghanistan	1999	745	19987071
Afghanistan	2000	2666	20595360
Brazil	1999	37737	172006362
Brazil	2000	80488	174504898
China	1999	212258	1272915272
China	2000	216766	128042583

variables

country	year	cases	population
Afghanistan	1999	745	19987071
Afghanistan	2000	2666	20595360
Brazil	1999	37737	172006362
Brazil	2000	80488	174504898
China	1999	212258	1272915272
China	2000	216766	128042583

observations

country	year	cases	population
Afghanistan	99	745	19987071
Afghanistan	00	2666	20595360
Brazil	99	37737	172006362
Brazil	00	80488	174504898
China	99	212258	1272915272
China	00	216766	128042583

values

“There are three interrelated rules which make a dataset tidy:

1. Each **variable** must have its own column.
2. Each **observation** must have its own row.
3. Each **value** must have its own cell.”

-Hadley Wickham, *R for Data Science*




Tidy data

```
#>   country      year  cases population
#>   <chr>      <int> <int>      <int>
#> 1 Afghanistan 1999     745   19987071
#> 2 Afghanistan 2000    2666   20595360
#> 3 Brazil       1999   37737  172006362
#> 4 Brazil       2000   80488  174504898
#> 5 China        1999  212258 1272915272
#> 6 China        2000  213766 1280428583
```

NOT tidy data


```
#>   country      `1999`      `2000`
#> * <chr>      <int>      <int>
#> 1 Afghanistan 19987071  20595360
#> 2 Brazil       172006362  174504898
#> 3 China        1272915272 1280428583
```

tidyr package helps with restructuring data




tidyr::gather(cases, "year", "n", 2:4)

Gather columns into rows.



tidyr::separate(storms, date, c("y", "m", "d"))

Separate one column into several.



tidyr::spread(pollution, size, amount)

Spread rows into columns.



tidyr::unite(data, col, ..., sep)

Unite several columns into one.

dplyr package helps with data transformation

Data Transformation with dplyr : : CHEAT SHEET




dplyr functions work with pipes and expect tidy data. In tidy data:



Summarise Cases

These apply **summary functions** to columns to create a new table of summary statistics. Summary functions take vectors as input and return one value (see back).




VARIATIONS

summarise_all() - Apply funs to every column.
summarise_at() - Apply funs to specific columns.
summarise_if() - Apply funs to all cols of one type.

Group Cases

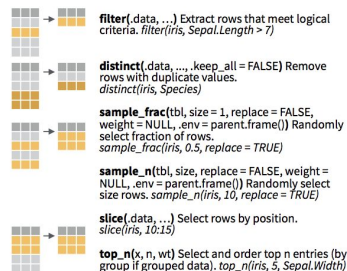
Use **group_by()** to create a "grouped" copy of a table. dplyr functions will manipulate each "group" separately and then combine the results.



Manipulate Cases

EXTRACT CASES

Row functions return a subset of rows as a new table.




Logical and boolean operators to use with filter()

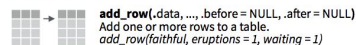
< **<=** **is.na()** **%in%** **|** **xor()**
> **>=** **!(is.na())** **!** **&**

See ?base::logic and ?Comparison for help.

ARRANGE CASES




ADD CASES



Manipulate Variables

EXTRACT VARIABLES

Column functions return a set of columns as a new vector or table.

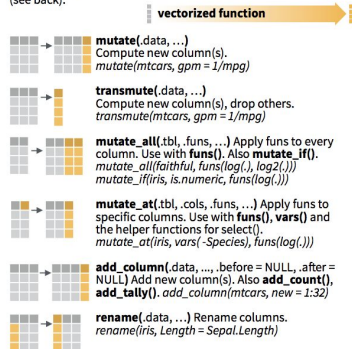


Use these helpers with **select()**, e.g. *select(iris, starts_with("Sepal"))*

contains(match) **num_range**(prefix, range) **%>%** mpg:cyl
ends_with(match) **one_of**(...) **%>%** g, Species
matches(match) **starts_with**(match)

MAKE NEW VARIABLES

These apply **vectorized functions** to columns. Vectorized funs take vectors as input and return vectors of the same length as output (see back).




dplyr sample functions


Manipulate Cases

EXTRACT CASES


Row functions return a subset of rows as a new table.




filter(.data, ...) Extract rows that meet logical criteria. *filter(iris, Sepal.Length > 7)*



distinct(.data, ..., .keep_all = FALSE) Remove rows with duplicate values. *distinct(iris, Species)*



sample_frac(tbl, size = 1, replace = FALSE, weight = NULL, .env = parent.frame()) Randomly select fraction of rows. *sample_frac(iris, 0.5, replace = TRUE)*



sample_n(tbl, size, replace = FALSE, weight = NULL, .env = parent.frame()) Randomly select size rows. *sample_n(iris, 10, replace = TRUE)*

slice(.data, ...) Select rows by position. *slice(iris, 10:15)*

top_n(x, n, wt) Select and order top n entries (by group if grouped data). *top_n(iris, 5, Sepal.Width)*

Logical and boolean operators to use with filter()

<	<=	is.na()	%in%		xor()
>	>=	!is.na()	!	&	

See **?base::logic** and **?Comparison** for help.

Manipulate Variables

EXTRACT VARIABLES

Column functions return a set of columns as a new vector or table.



pull(.data, var = -1) Extract column values as a vector. Choose by name or index. *pull(iris, Sepal.Length)*



select(.data, ...) Extract columns as a table. Also **select_if()**. *select(iris, Sepal.Length, Species)*

Use these helpers with select (),
e.g. *select(iris, starts_with("Sepal"))*

contains (match)	num_range (prefix, range)	;, e.g. mpg:cyl
ends_with (match)	one_of (...)	-, e.g. -Species
matches (match)	starts_with (match)	

MAKE NEW VARIABLES

These apply **vectorized functions** to columns. Vectorized funs take vectors as input and return vectors of the same length as output (see back).

vectorized function



mutate(.data, ...) Compute new column(s). *mutate(mtcars, gpm = 1/mpg)*



transmute(.data, ...) Compute new column(s), drop others. *transmute(mtcars, gpm = 1/mpg)*

stringr package helps with string manipulation

String manipulation with stringr : : CHEAT SHEET



The **stringr** package provides a set of internally consistent tools for working with character strings, i.e. sequences of characters surrounded by quotation marks.

Detect Matches



str_detect(string, pattern) Detect the presence of a pattern match in a string.
str_detect(fruit, "a")



str_which(string, pattern) Find the indexes of strings that contain a pattern match.
str_which(fruit, "a")



str_count(string, pattern) Count the number of matches in a string.
str_count(fruit, "a")



str_locate(string, pattern) Locate the positions of pattern matches in a string. Also **str_locate_all**.
str_locate(fruit, "a")

Subset Strings



str_sub(string, start = 1L, end = -1L) Extract substrings from a character vector.
str_sub(fruit, 1, 3); str_sub(fruit, -2)



str_subset(string, pattern) Return only the strings that contain a pattern match.
str_subset(fruit, "b")



str_extract(string, pattern) Return the first pattern match found in each string, as a vector. Also **str_extract_all** to return every pattern match.
str_extract(fruit, "[aeiou]")



str_match(string, pattern) Return the first pattern match found in each string, as a matrix with a column for each () group in pattern. Also **str_match_all**.
str_match(sentences, "[a]the ([^]+)")

Manage Lengths



str_length(string) The width of strings (i.e. number of code points, which generally equals the number of characters). *str_length(fruit)*



str_pad(string, width, side = c("left", "right", "both"), pad = " ") Pad strings to constant width. *str_pad(fruit, 17)*



str_trunc(string, width, side = c("right", "left", "center"), ellipsis = "...") Truncate the width of strings, replacing content with ellipsis.
str_trunc(fruit, 3)



str_trim(string, side = c("both", "left", "right")) Trim whitespace from the start and/or end of a string. *str_trim(fruit)*

Mutate Strings



str_sub() <- value. Replace substrings by identifying the substrings with **str_sub()** and assigning into the results.
str_sub(fruit, 1, 3) <- "str"



str_replace(string, pattern, replacement) Replace the first matched pattern in each string. *str_replace(fruit, "a", "-")*



str_replace_all(string, pattern, replacement) Replace all matched patterns in each string. *str_replace_all(fruit, "a", "-")*



str_to_lower(string, locale = "en") Convert strings to lower case. *str_to_lower(sentences)*



str_to_upper(string, locale = "en") Convert strings to upper case. *str_to_upper(sentences)*



str_to_title(string, locale = "en") Convert strings to title case. *str_to_title(sentences)*

Join and Split



str_c(..., sep = "", collapse = NULL) Join multiple strings into a single string.
str_c(letters, LETTERS)



str_c(..., sep = "", collapse = NULL) Collapse a vector of strings into a single string.
str_c(letters, collapse = "")



str_dup(string, times) Repeat strings times times. *str_dup(fruit, times = 2)*



str_split_fixed(string, pattern, n) Split a vector of strings into a matrix of substrings (splitting at occurrences of a pattern match). Also **str_split** to return a list of substrings.
str_split_fixed(fruit, "", n=2)



str_glue(..., sep = "", envir = parent.frame()) Create a string from strings and (expressions) to evaluate. *str_glue("Pi is {pi}")*



str_glue_data(x, ..., sep = "", envir = parent.frame(), na = "NA") Use a data frame, list, or environment to create a string from strings and (expressions) to evaluate.
str_glue_data(mtcars, "(rownames(mtcars)) has {hp} hp")

Order Strings



str_order(x, decreasing = FALSE, na_last = TRUE, locale = "en", numeric = FALSE, ...) Return the vector of indexes that sorts a character vector. *x[str_order(x)]*



str_sort(x, decreasing = FALSE, na_last = TRUE, locale = "en", numeric = FALSE, ...) Sort a character vector.
str_sort(x)

Helpers



str_conv(string, encoding) Override the encoding of a string. *str_conv(fruit, "ISO-8859-1")*

str_view(string, pattern, match = NA) View HTML rendering of first regex match in each string. *str_view(fruit, "[aeiou]")*

str_view_all(string, pattern, match = NA) View HTML rendering of all regex matches.
str_view_all(fruit, "[aeiou]")

str_wrap(string, width = 80, indent = 0, exdent = 0) Wrap strings into nicely formatted paragraphs. *str_wrap(sentences, 20)*

stringr sample functions

Mutate Strings



str_sub() <- value. Replace substrings by identifying the substrings with `str_sub()` and assigning into the results.
`str_sub(fruit, 1, 3) <- "str"`



str_replace()(string, **pattern**, replacement) Replace the first matched pattern in each string. `str_replace(fruit, "a", "-")`



str_replace_all()(string, **pattern**, replacement) Replace all matched patterns in each string. `str_replace_all(fruit, "a", "-")`

A STRING
↓
a string

str_to_lower()(string, locale = "en")¹ Convert strings to lower case.
`str_to_lower(sentences)`

a string
↓
A STRING

str_to_upper()(string, locale = "en")¹ Convert strings to upper case.
`str_to_upper(sentences)`

a string
↓
A String

str_to_title()(string, locale = "en")¹ Convert strings to title case. `str_to_title(sentences)`

Join and Split



str_c(..., sep = "", collapse = NULL) Join multiple strings into a single string.
`str_c(letters, LETTERS)`



str_c(..., sep = "", collapse = NULL) Collapse a vector of strings into a single string.
`str_c(letters, collapse = "")`



str_dup()(string, times) Repeat strings times times. `str_dup(fruit, times = 2)`



str_split_fixed()(string, **pattern**, n) Split a vector of strings into a matrix of substrings (splitting at occurrences of a pattern match). Also **str_split** to return a list of substrings.
`str_split_fixed(fruit, " ", n=2)`



str_glue(..., .sep = "", .envir = parent.frame()) Create a string from strings and {expressions} to evaluate. `str_glue("Pi is {pi}")`



str_glue_data(x, ..., .sep = "", .envir = parent.frame(), .na = "NA") Use a data frame, list, or environment to create a string from strings and {expressions} to evaluate.
`str_glue_data(mtcars, "{rownames(mtcars)} has {hp} hp")`

Helpful cheatsheets for data cleaning with R

[Data wrangling: tidyr & dplyr](#)

[Data Transformation with dplyr](#) (in-depth focus on dplyr functions)

[String manipulation: stringr](#)

[All other R cheatsheets](#) (from RStudio)